

## **Module II**

# **Concurrency Control Techniques**

# Introduction

- Concurrency control protocols
  - Set of rules to guarantee serializability
    1. Two-phase locking protocols
      - Lock data items to prevent concurrent access
    2. Timestamp
      - Unique identifier for each transaction
    3. Multiversion currency control protocols
      - Use multiple versions of a data item
- Validation or certification of a transaction

# 21.1 Two-Phase Locking Techniques for Concurrency Control

- Lock
  - Variable associated with a data item describing status for operations that can be applied
  - One lock for each item in the database
- Binary locks
  - Two states (values)
    - Locked (1)
      - Item cannot be accessed
    - Unlocked (0)
      - Item can be accessed when requested

# Two-Phase Locking Techniques for Concurrency Control (cont'd.)

- Transaction requests access by issuing a `lock_item(X)` operation

```
lock_item(X):  
  B:  if LOCK(X) = 0          (*item is unlocked*)  
        then LOCK(X) ← 1      (*lock the item*)  
      else  
        begin  
          wait (until LOCK(X) = 0  
                and the lock manager wakes up the transaction);  
          go to B  
        end;  
unlock_item(X):  
  LOCK(X) ← 0;                (* unlock the item *)  
  if any transactions are waiting  
    then wakeup one of the waiting transactions;
```

Figure 21.1 Lock and unlock operations for binary locks

# Two-Phase Locking Techniques for Concurrency Control (cont'd.)

- Lock table specifies items that have locks
- Lock manager subsystem
  - Keeps track of and controls access to locks
  - Rules enforced by lock manager module
- At most one transaction can hold the lock on an item at a given time
- Binary locking too restrictive for database items

# Two-Phase Locking Techniques for Concurrency Control (cont'd.)

- Shared/exclusive or read/write locks
  - Read operations on the same item are not conflicting
  - Must have exclusive lock to write
  - Three locking operations
    - `read_lock(X)`
    - `write_lock(X)`
    - `unlock(X)`

Figure 21.2 Locking and unlocking operations for two-mode (read/write, or shared/exclusive) locks

```
read_lock(X):
B:  if LOCK(X) = "unlocked"
    then begin LOCK(X) ← "read-locked";
        no_of_reads(X) ← 1
    end
    else if LOCK(X) = "read-locked"
        then no_of_reads(X) ← no_of_reads(X) + 1
    else begin
        wait (until LOCK(X) = "unlocked"
            and the lock manager wakes up the transaction);
        go to B
    end;

write_lock(X):
B:  if LOCK(X) = "unlocked"
    then LOCK(X) ← "write-locked"
    else begin
        wait (until LOCK(X) = "unlocked"
            and the lock manager wakes up the transaction);
        go to B
    end;

unlock (X):
    if LOCK(X) = "write-locked"
        then begin LOCK(X) ← "unlocked";
            wakeup one of the waiting transactions, if any
        end
    else if LOCK(X) = "read-locked"
        then begin
            no_of_reads(X) ← no_of_reads(X) - 1;
            if no_of_reads(X) = 0
                then begin LOCK(X) = "unlocked";
                    wakeup one of the waiting transactions, if any
                end
        end
    end;
```

# Two-Phase Locking Techniques for Concurrency Control (cont'd.)

- Lock conversion
  - Transaction that already holds a lock allowed to convert the lock from one state to another
- Upgrading
  - Issue a read\_lock operation then a write\_lock operation
- Downgrading
  - Issue a read\_lock operation after a write\_lock operation



# Guaranteeing Serializability by Two-Phase Locking

- Two-phase locking protocol
  - All locking operations precede the first unlock operation in the transaction
  - Phases
    - Expanding (growing) phase
      - New locks can be acquired but none can be released
      - Lock conversion upgrades must be done during this phase
    - Shrinking phase
      - Existing locks can be released but none can be acquired
      - Downgrades must be done during this phase

Figure 21.3 Transactions that do not obey two-phase locking (a) Two transactions  $T_1$  and  $T_2$  (b) Results of possible serial schedules of  $T_1$  and  $T_2$  (c) A nonserializable schedule  $S$  that uses locks

(a)	$T_1$	$T_2$
	<pre> read_lock(Y); read_item(Y); unlock(Y); write_lock(X); read_item(X); X := X + Y; write_item(X); unlock(X); </pre>	<pre> read_lock(X); read_item(X); unlock(X); write_lock(Y); read_item(Y); Y := X + Y; write_item(Y); unlock(Y); </pre>

(b) Initial values:  $X=20$ ,  $Y=30$

Result serial schedule  $T_1$   
followed by  $T_2$ :  $X=50, Y=80$

Result of serial schedule  $T_2$  followed by  $T_1$ :  $X=70, Y=50$

(c)

	$T_1$	$T_2$
Time ↓	read_lock(Y); read_item(Y); unlock(Y);	read_lock(X); read_item(X); unlock(X); write_lock(Y); read_item(Y); $Y := X + Y$ ; write_item(Y); unlock(Y);
	write_lock(X); read_item(X); $X := X + Y$ ; write_item(X); unlock(X);	

# Guaranteeing Serializability by Two-Phase Locking

- If every transaction in a schedule follows the two-phase locking protocol, schedule guaranteed to be serializable
- Two-phase locking may limit the amount of concurrency that can occur in a schedule
- Some serializable schedules will be prohibited by two-phase locking protocol

# Variations of Two-Phase Locking

- Basic 2PL
  - Technique described on previous slides
- Strict 2PL
  - Transaction does not release exclusive locks until after it commits or aborts
- Conservative (static) 2PL
  - Requires a transaction to lock all the items it accesses before the transaction begins
    - Predeclare read-set and write-set
  - Deadlock-free protocol

# Variations of Two-Phase Locking (cont'd.)

- Rigorous 2PL
  - Transaction does not release any locks until after it commits or aborts
- Concurrency control subsystem responsible for generating read\_lock and write\_lock requests
- Locking generally considered to have high overhead

# Dealing with Deadlock and Starvation

## ■ Deadlock

- Occurs when each transaction  $T$  in a set is waiting for some item locked by some other transaction  $T'$
- Both transactions stuck in a waiting queue

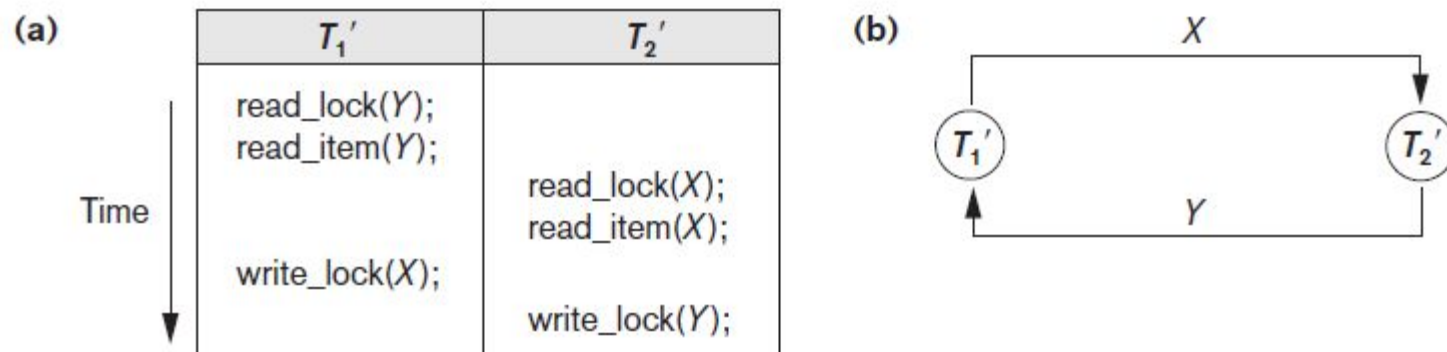


Figure 21.5 Illustrating the deadlock problem (a) A partial schedule of  $T_1'$  and  $T_2'$  that is in a state of deadlock (b) A wait-for graph for the partial schedule in (a)

# Dealing with Deadlock and Starvation (cont'd.)

- Deadlock prevention protocols
  - Every transaction locks all items it needs in advance
  - Ordering all items in the database
    - Transaction that needs several items will lock them in that order
  - Both approaches impractical
- Protocols based on a timestamp
  - Wait-die
  - Wound-wait

# Dealing with Deadlock and Starvation (cont'd.)

- No waiting algorithm
  - If transaction unable to obtain a lock, immediately aborted and restarted later
- Cautious waiting algorithm
  - Deadlock-free
- Deadlock detection
  - System checks to see if a state of deadlock exists
  - Wait-for graph



# Dealing with Deadlock and Starvation (cont'd.)

- Victim selection
  - Deciding which transaction to abort in case of deadlock
- Timeouts
  - If system waits longer than a predefined time, it aborts the transaction
- Starvation
  - Occurs if a transaction cannot proceed for an indefinite period of time while other transactions continue normally
  - Solution: first-come-first-served queue

# 21.2 Concurrency Control Based on Timestamp Ordering

## ■ Timestamp

- Ordering among the transactions is determined in advance based on their time stamp
- The time at which transaction enters into system for execution (i.e. Transaction start time)
- Unique identifier assigned by the DBMS to identify a transaction
- Assigned in the order submitted
- If any transaction  $T_j$  enters after  $T_i$ , the relation between time stamp will be  $TS(T_i) < TS(T_j)$  which means that the resulting schedule must be equivalent to a serial schedule  $T_i \rightarrow T_j$

# Concurrency Control Based on Timestamp Ordering (cont'd.)

- Generating timestamps
  - Counter incremented each time its value is assigned to a transaction
  - Current date/time value of the system clock
    - Ensure no two timestamps are generated during the same tick of the clock
- General approach
  - Enforce equivalent serial order on the transactions based on their timestamps

# Concurrency Control Based on Timestamp Ordering (cont'd.)

- Concurrency control techniques based on timestamps do not use locks
  - Deadlocks cannot occur
- Timestamp ordering (TO)
  - Allows interleaving of transaction operations
  - Must ensure timestamp order is followed for each pair of conflicting operations
- Each database item assigned two timestamp values
  - $read\_TS(X)$
  - $write\_TS(X)$

# Concurrency Control Based on Timestamp Ordering (cont'd.)

## ■ Basic TO algorithm

- If conflicting operations detected, later operation rejected by aborting transaction that issued it
- Schedules produced guaranteed to be conflict serializable
- Starvation may occur

## ■ Strict TO algorithm

- Ensures schedules are both strict and conflict serializable

# Concurrency Control Based on Timestamp Ordering (cont'd.)

- Thomas's write rule
  - Modification of basic TO algorithm
  - Does not enforce conflict serializability
  - Rejects fewer write operations by modifying checks for write\_item(X) operation
  - Ignoring outdated writes is called the Thomas write rule
  - Ex- R1(A)W2(A)W1(A)

# 21.3 Multiversion Concurrency Control Techniques

- Several versions of an item are kept by a system
- Some read operations that would be rejected in other techniques can be accepted by reading an older version of the item
  - Maintains serializability
- More storage is needed
- Multiversion concurrency control scheme types
  - Based on two-phase locking
  - Based on timestamp ordering
  - Validation and snapshot isolation techniques

# Multiversion Concurrency Control Techniques (cont'd.)

- Multiversion two-phase locking using certify locks
  - Three locking modes: read, write, and certify

(a)		Read	Write
		<hr/>	
	Read	Yes	No
	Write	No	No

(b)		Read	Write	Certify
		<hr/>		
	Read	Yes	Yes	No
	Write	Yes	No	No
	Certify	No	No	No

Figure 21.6 Lock compatibility tables (a) Lock compatibility table for read/write locking scheme (b) Lock compatibility table for read/write/certify locking scheme



# Multiversion Concurrency Control Techniques (cont'd.)

- Allow other transaction T2 to read an item X while T1 holds a write lock on X.
- Maintain two version of for each item: **committed** and **local** version.
- Committed version must always be written by some committed transaction.
- Local version created when a transaction obtain a write lock on X.
- Obtain certify lock before commit.

# Multiversion Concurrency Control Techniques (cont'd.)

- Several version of data item X such as X1, X2, X3,...,Xk are maintained
- Multiversion technique based on timestamp ordering
  - Two timestamps associated with each version are kept
    - $read\_TS(X_i)$ -largest of all the timestamps of a transaction that have successfully read version  $X_i$
    - $write\_TS(X_i)$ - Timestamp of the transaction that wrote the value of version  $X_i$

- Ensure serializability based on following rules:

2. If transaction  $T$  issues a  $\text{write\_item}(X)$  operation, and version  $i$  of  $X$  has the highest  $\text{write\_TS}(Xi)$  of all versions of  $X$  that is also *less than or equal to*  $\text{TS}(T)$ , and  $\text{read\_TS}(Xi) > \text{TS}(T)$ , then abort and roll back transaction  $T$ ; otherwise, create a new version  $Xj$  of  $X$  with  $\text{read\_TS}(Xj) = \text{write\_TS}(Xj) = \text{TS}(T)$ .
3. If transaction  $T$  issues a  $\text{read\_item}(X)$  operation, find the version  $i$  of  $X$  that has the highest  $\text{write\_TS}(Xi)$  of all versions of  $X$  that is also *less than or equal to*  $\text{TS}(T)$ ; then return the value of  $Xi$  to transaction  $T$ , and set the value of  $\text{read\_TS}(Xi)$  to the larger of  $\text{TS}(T)$  and the current  $\text{read\_TS}(Xi)$ .

TS(T1) = 4      TS (T2) = 2      TS(T3) = 8      T4 TS(T4)= 6

W(A)

W(A)

W(A)

A1 (10, 2, R.Ts)

A2 (20, 4, R.TS)

A(30, 6, R.TS)

R\_4(A2) Allowed

W\_4(A3) not allowed. Abort the transaction

W\_4(A4) = TS(T4) allowed and a new version is created.

# 21.4 Validation (Optimistic) Techniques and Snapshot Isolation Concurrency Control

- Optimistic techniques

- Also called validation or certification techniques
- No checking is done while the transaction is executing
- Updates not applied directly to the database until finished transaction is validated
  - All updates applied to local copies of data items
- Validation phase checks whether any of transaction's updates violate serializability
  - Transaction committed or aborted based on result

# Concurrency Control Based on Snapshot Isolation

- Transaction sees data items based on committed values of the items in the database snapshot
  - Does not see updates that occur after transaction starts
- Read operations do not require read locks
  - Write operations require write locks
- Temporary version store keeps track of older versions of updated items
- Variation: serializable snapshot isolation (SSI)

# 21.8 Summary

- Concurrency control techniques
  - Two-phase locking
  - Timestamp-based ordering
  - Multiversion protocols
  - Snapshot isolation