

Transactions

by

Dr. Pratik Roy

Definition

- ❑ A transaction is a unit of program execution that accesses and possibly updates various data items.
- ❑ A transaction is an executing program that forms a logical unit of database processing.
- ❑ **Example:** Let T is a transaction that transfers Rs. 50/- from account A to account B:

```
T: read(A);  
    $A := A - 50;$   
   write(A);  
   read(B);  
    $B := B + 50;$   
   write(B).
```

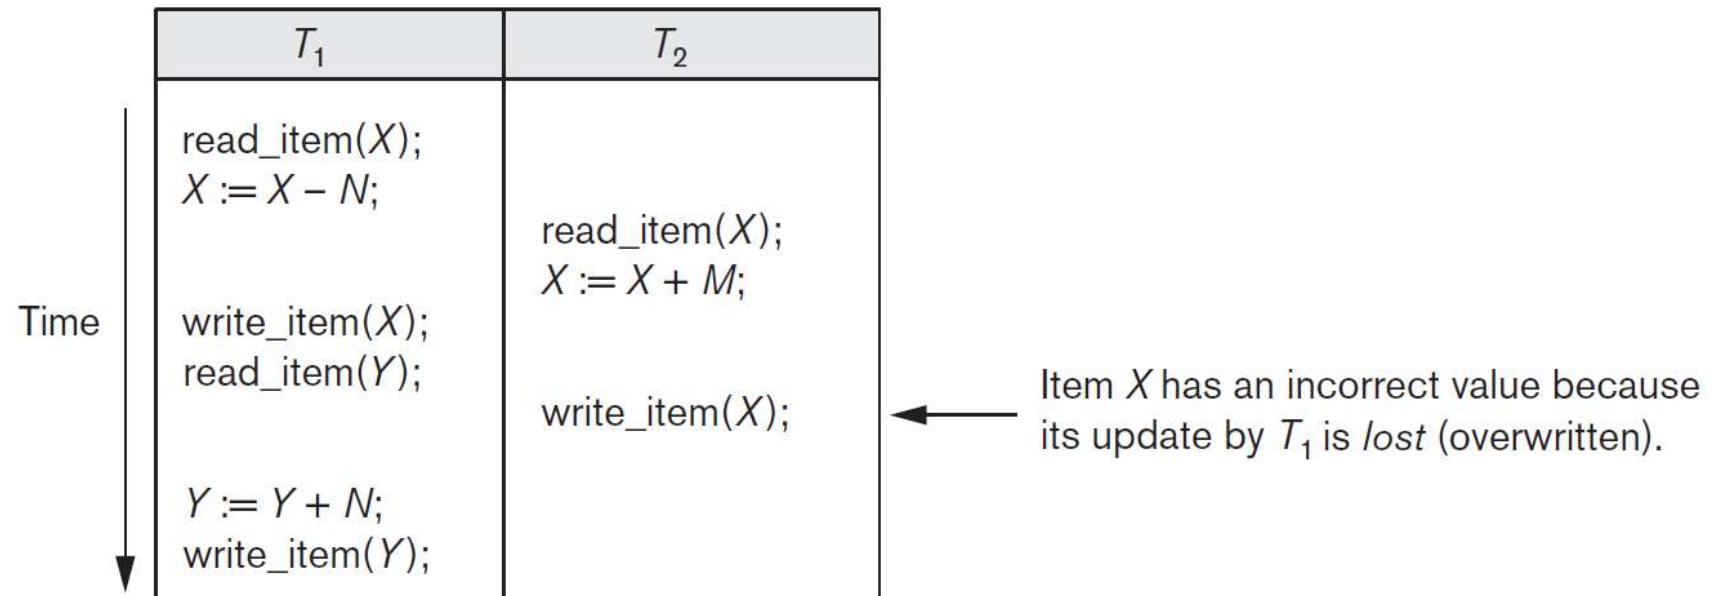
Concurrency Control

- ❑ Transactions submitted by various users may execute concurrently.
 - Access and update the same database items
 - Some form of concurrency control is needed
- ❑ Two main issues to deal with:
 - Failures of various kinds, such as hardware failures and system crashes
 - Concurrent execution of multiple transactions

Why Concurrency Control Is Needed

❑ Lost Update Problem

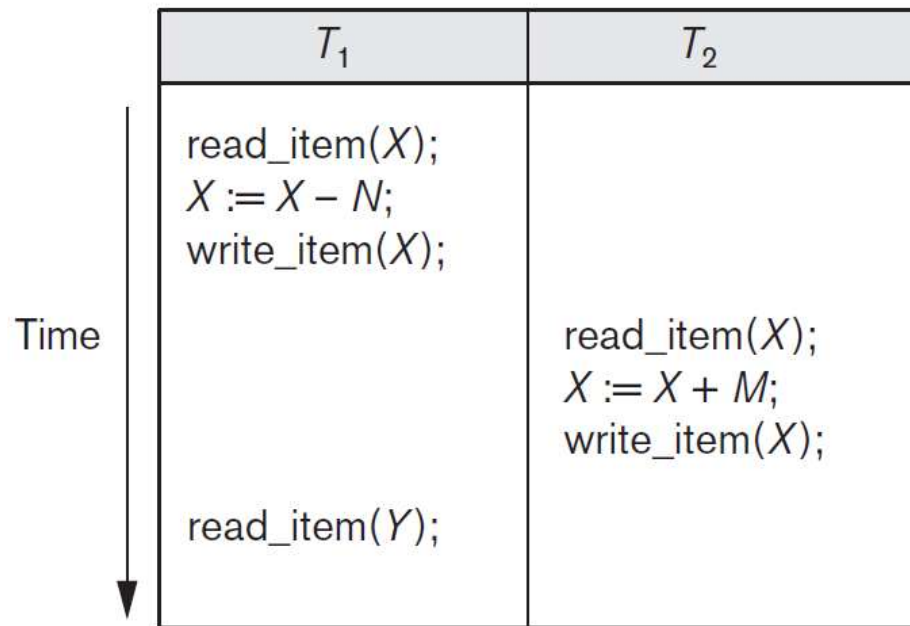
- Occurs when two transactions that access the same database items have their operations interleaved in a way that makes the value of some database items incorrect.



Why Concurrency Control Is Needed

□ Temporary Update (or Dirty Read) Problem

- Occurs when one transaction updates a database item and then the transaction fails for some reason.
- Meanwhile, the updated item is accessed (read) by another transaction before it is changed back to its original value.



Transaction T_1 fails and must change the value of X back to its old value; meanwhile T_2 has read the *temporary* incorrect value of X .

T_1 updates item X and then fails before completion

Why Concurrency Control Is Needed

❑ Incorrect Summary Problem

- If one transaction is calculating an aggregate summary function on a number of database items while other transactions are updating some of these items, the aggregate function may calculate some values before they are updated and others after they are updated.

T_1	T_3
<pre>read_item(X); X := X - N; write_item(X); read_item(Y); Y := Y + N; write_item(Y);</pre>	<pre>sum := 0; read_item(A); sum := sum + A; . . . read_item(X); sum := sum + X; read_item(Y); sum := sum + Y;</pre>

← T_3 reads X after N is subtracted and reads Y before N is added; a wrong summary is the result (off by N).

Why Concurrency Control Is Needed

□ Unrepeatable Read Problem

- Transaction T reads the same item twice and the item is changed by another transaction T' between the two reads.
- T receives different values for its two reads of the same item.

Required Properties of a Transaction

❑ Consider a transaction to transfer Rs. 50/- from account A to account B:

read(A)

$A := A - 50$

write(A)

read(B)

$B := B + 50$

write(B)

❑ Atomicity requirement

- If the transaction fails after step 3 and before step 6, money will be “lost” leading to an inconsistent database state
 - Failure could be due to software or hardware
- The system should ensure that updates of a partially executed transaction are not reflected in the database

Required Properties of a Transaction

□ Durability requirement

- Once the user has been notified that the transaction has completed (i.e., the transfer of the Rs. 50/- has taken place), the updates to the database by the transaction must persist even if there are software or hardware failures.

□ Consistency requirement in above example:

- Sum of A and B is unchanged by the execution of the transaction

□ In general, consistency requirements include

- Explicitly specified integrity constraints such as primary keys and foreign keys
- Implicit integrity constraints
 - e.g., sum of balances of all accounts, minus sum of loan amounts must equal value of cash-in-hand

□ A transaction, when starting to execute, must see a consistent database.

□ During transaction execution the database may be temporarily inconsistent.

□ When the transaction completes successfully the database must be consistent

- Erroneous transaction logic can lead to inconsistency

Required Properties of a Transaction

- ❑ **Isolation requirement:** If between steps 3 and 6 (of the fund transfer transaction), another transaction T2 is allowed to access the partially updated database, it will see an inconsistent database (the sum $A + B$ will be less than it should be).

T1
read(A)
 $A := A - 50$
write(A)

read(B)
 $B := B + 50$
write(B)

T2

read(A), read(B), print(A+B)

- ❑ Isolation can be ensured trivially by running transactions **serially**
 - That is, one after the other.
- ❑ Executing multiple transactions concurrently has significant benefits.

ACID Properties

❑ To preserve the integrity of data the database system must ensure:

❑ **Atomicity** - Either all operations of the transaction are properly reflected in the database or none are.

- Transaction performed in its entirety or not at all

❑ **Consistency** - Execution of a transaction in isolation preserves the consistency of the database.

- Takes database from one consistent state to another

❑ **Isolation** - Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.

- For every pair of transactions T_i and T_j , it appears to T_i that either T_j finished execution before T_i started, or T_j started execution after T_i finished.

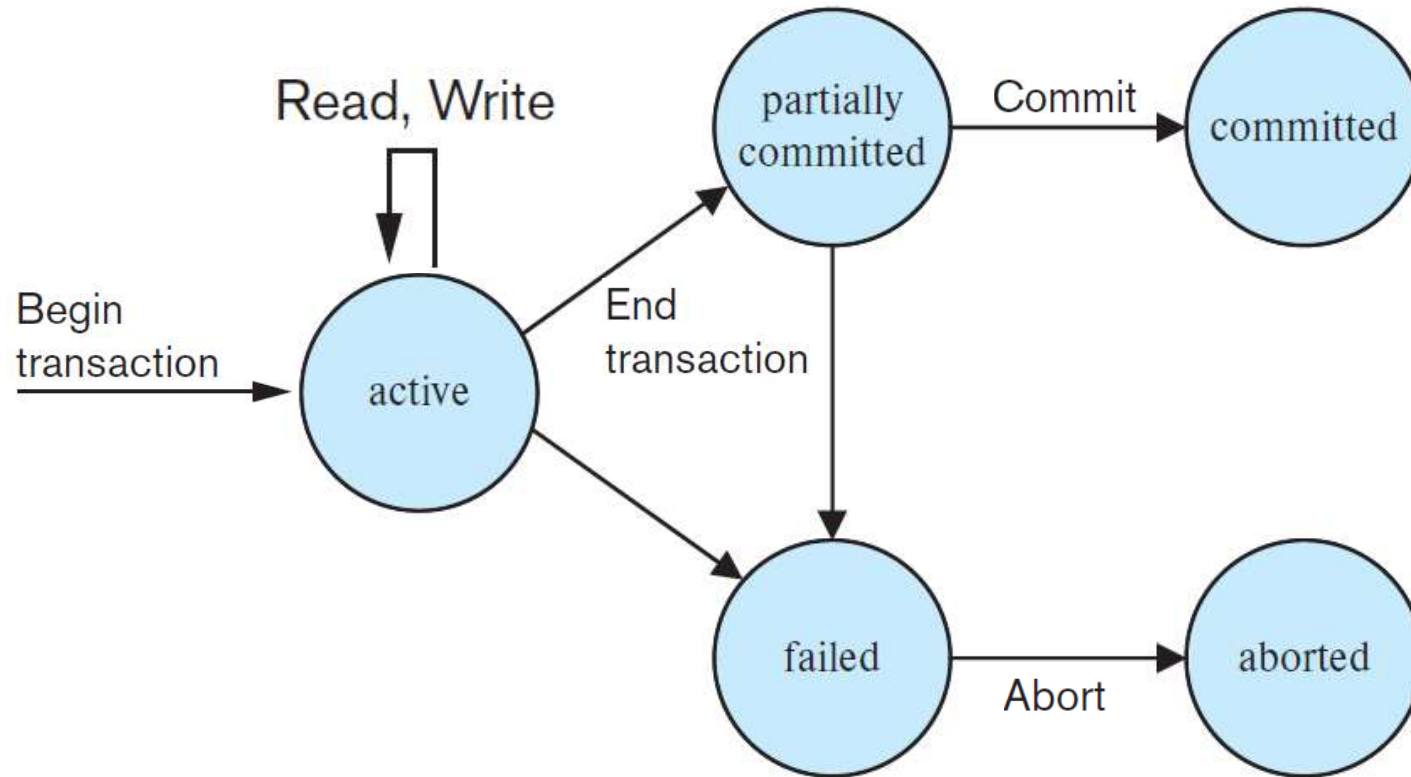
ACID Properties

- ❑ **Durability** - After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.

Transaction State

- ❑ **Active:** The initial state; the transaction stays in this state while it is executing.
- ❑ **Partially committed:** After the final statement has been executed.
- ❑ **Failed:** After the discovery that normal execution can no longer proceed.
- ❑ **Aborted:** After the transaction has been rolled back and the database has been restored to its state prior to the start of the transaction. Two options after it has been aborted:
 - Restart the transaction
 - can be done only if no internal logical error
 - Kill the transaction
- ❑ **Committed:** After successful completion.

State diagram of Transaction



Concurrent Executions

❑ Multiple transactions are allowed to run concurrently in the system.

Advantages are:

- **Increased processor and disk utilization**, leading to better transaction throughput
 - E.g. one transaction can be using the CPU while another is reading from or writing to the disk
- **Reduced average response time** for transactions: short transactions need not wait behind long ones.

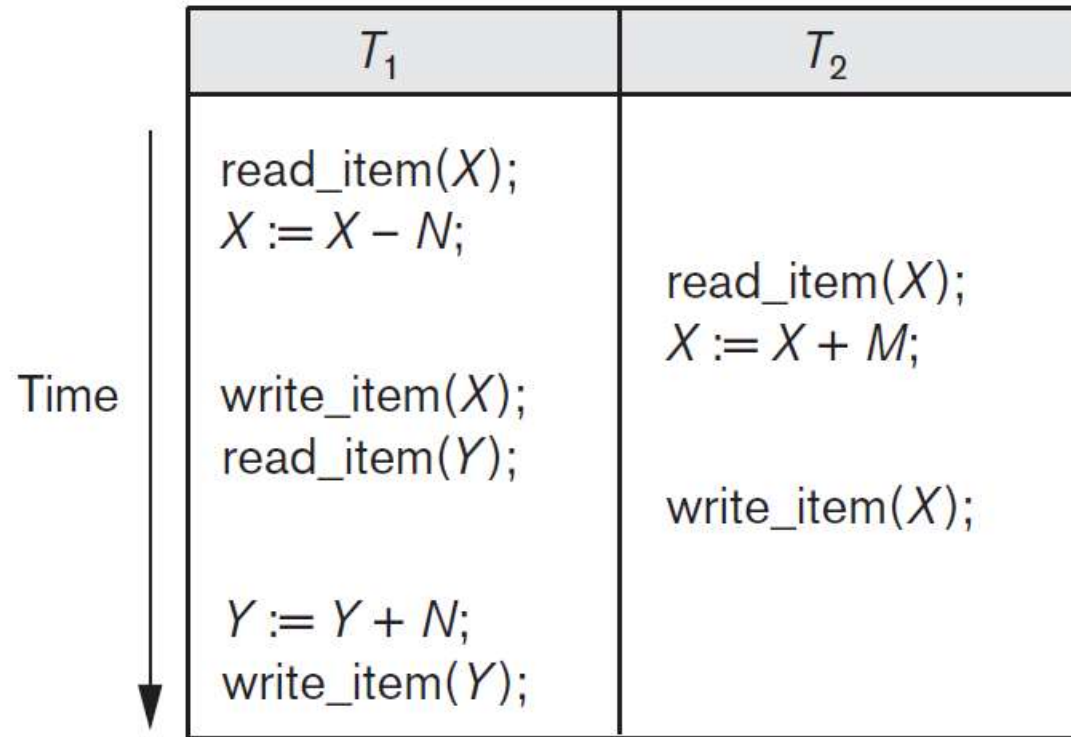
❑ **Concurrency control schemes** – mechanisms to achieve isolation

- To control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database.

Schedules (Histories) of Transactions

- ❑ A sequence of instructions that specify the chronological order in which instructions are executed in the system.
- ❑ Order of execution of operations from all transactions.
- ❑ A schedule for a set of transactions must consist of all instructions of those transactions.
- ❑ Operations from different transactions can be interleaved in the schedule.
- ❑ Must preserve the order in which the instructions appear in each individual transaction.
- ❑ Total ordering of operations in a schedule
 - For any two operations in the schedule, one must occur before the other.

Example of Schedule

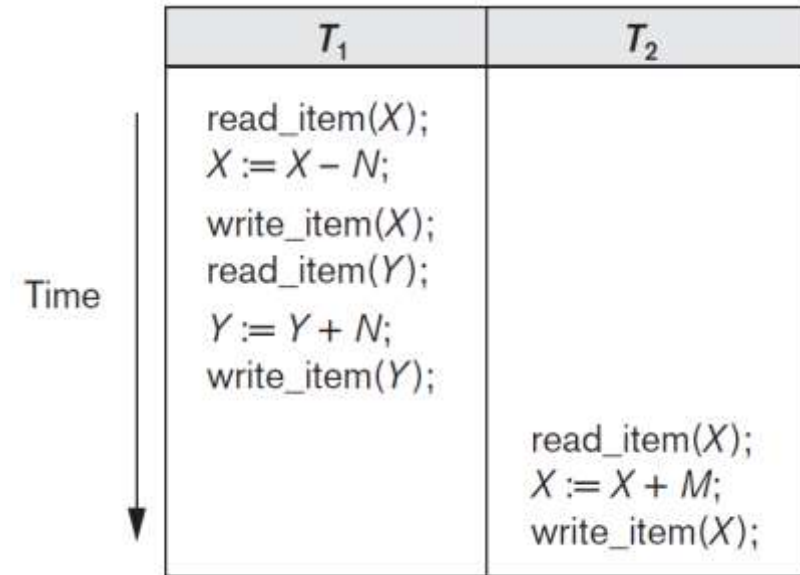


$S_a: r_1(X); r_2(X); w_1(X); r_1(Y); w_2(X); w_1(Y);$

Serial Schedule

- Each serial schedule consists of a sequence of instructions from various transactions, where the instructions belonging to one single transaction appear together in that schedule.

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit



T_1 is followed by T_2

Serial Schedule

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Time

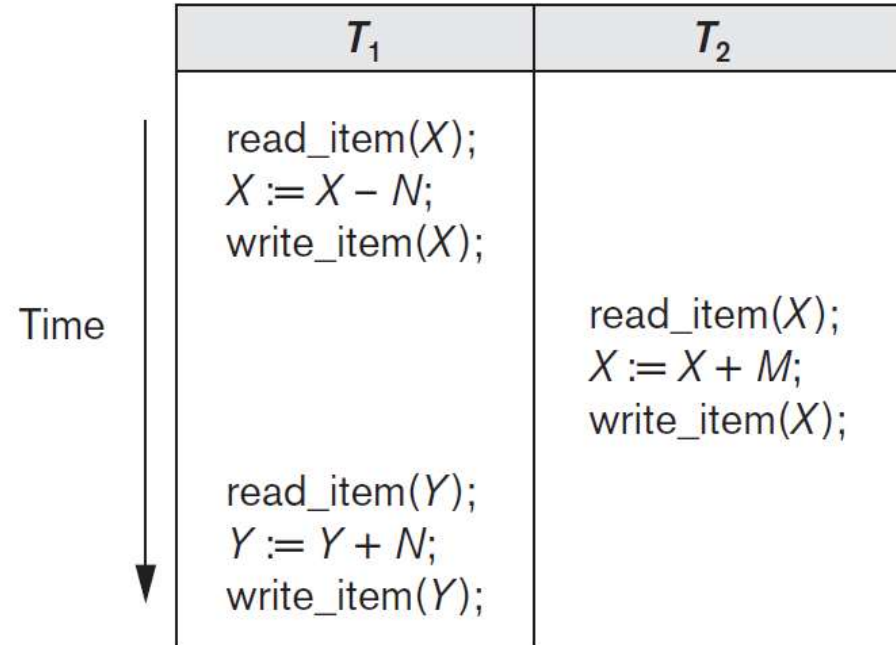


T_1	T_2
read_item(X); $X := X - N$; write_item(X); read_item(Y); $Y := Y + N$; write_item(Y);	read_item(X); $X := X + M$; write_item(X);

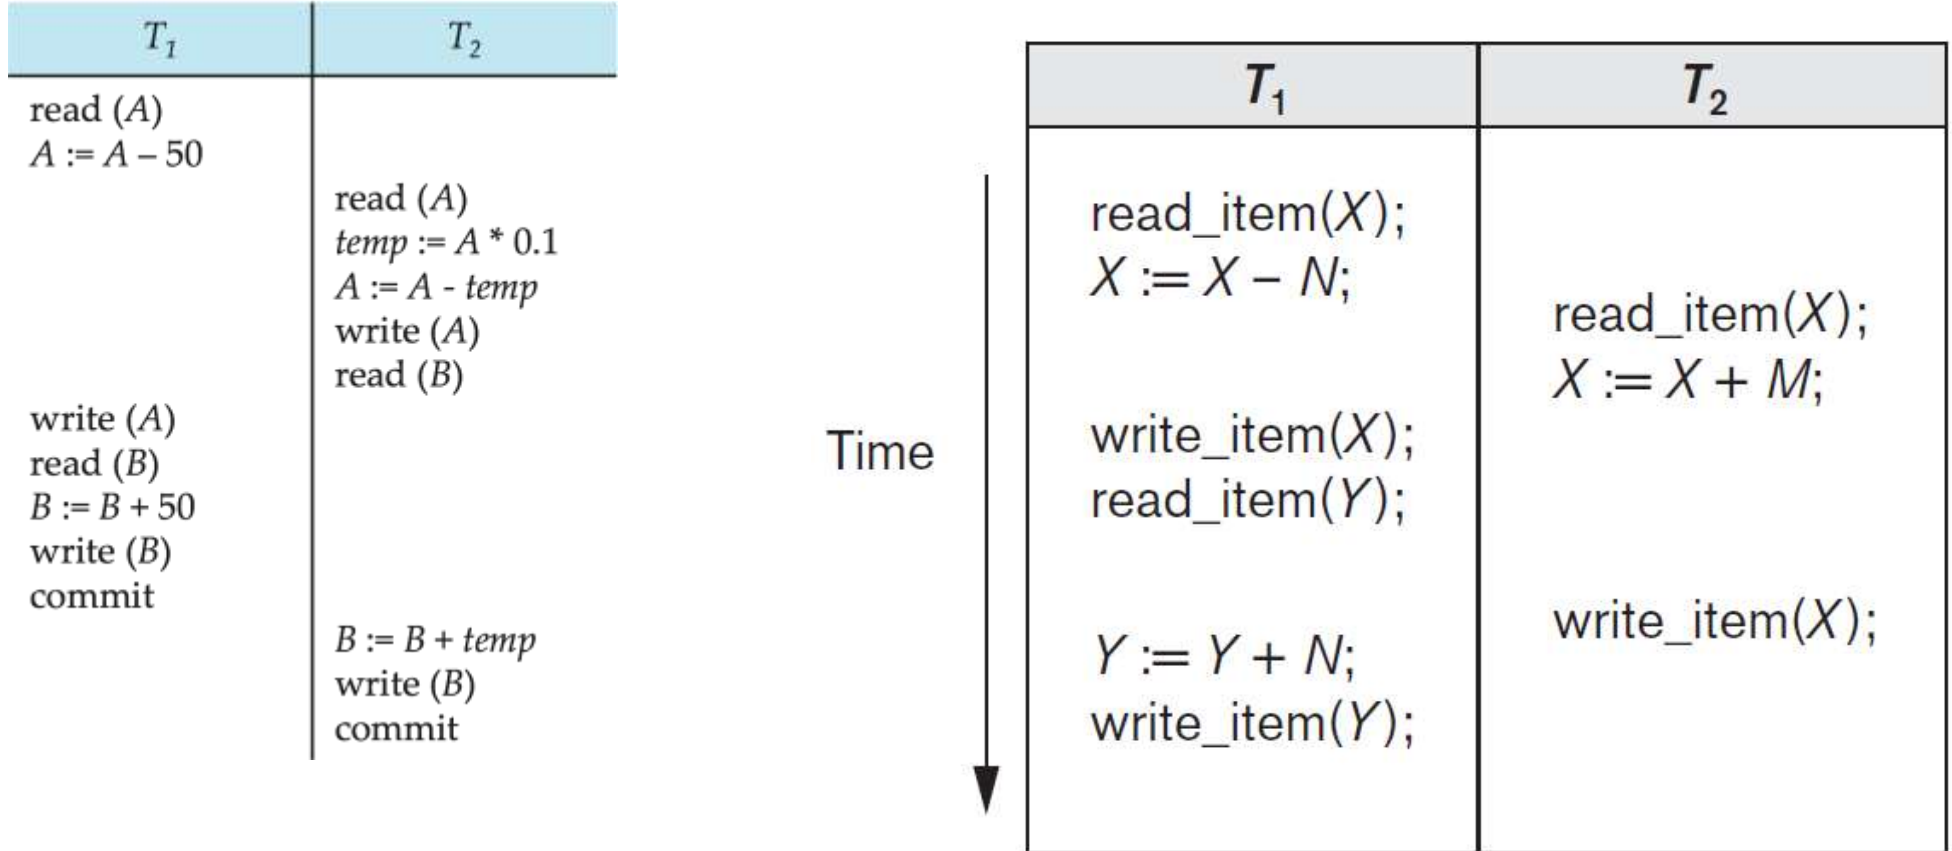
T_2 is followed by T_1

Nonserial schedule

T_1	T_2
read (A) $A := A - 50$ write (A)	
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	
	read (B) $B := B + temp$ write (B) commit



Nonserial schedule



Give erroneous result (inconsistent state)

Serializability

- ❑ **Basic Assumption** – Each transaction preserves database consistency.
- ❑ Serial execution of a set of transactions preserves database consistency.
- ❑ **Problem with serial schedules**
 - Limit concurrency by prohibiting interleaving of operations
 - Unacceptable in practice
 - **Solution:** determine which schedules are equivalent to a serial schedule and allow those to occur
- ❑ A (possibly concurrent) schedule is **serializable** if it is equivalent to a serial schedule.
- ❑ Serializable schedule of n transactions
 - Equivalent to some serial schedule of same n transactions
- ❑ Different forms of schedule equivalence give rise to the notions of:
 - **Conflict serializability**
 - **View serializability**

Conflicting Instructions

- ❑ Let us consider a schedule S in which there are two consecutive instructions, I and J , of transactions T_i and T_j , respectively ($i \neq j$).
- ❑ If I and J refer to different data items, then we can swap I and J without affecting the results of any instruction in the schedule.
- ❑ If I and J refer to the same data item Q , then the order of the two steps may matter.
 1. **$I = \text{read}(Q)$, $J = \text{read}(Q)$** . The order of I and J does not matter, since the same value of Q is read by T_i and T_j , regardless of the order.
 2. **$I = \text{read}(Q)$, $J = \text{write}(Q)$** . If I comes before J , then T_i does not read the value of Q that is written by T_j in instruction J . If J comes before I , then T_i reads the value of Q that is written by T_j . Thus, the order of I and J matters.
 3. **$I = \text{write}(Q)$, $J = \text{read}(Q)$** . The order of I and J matters.

Conflicting Instructions

4. $I = \text{write}(Q)$, $J = \text{write}(Q)$. Since both instructions are write operations, the order of these instructions does not affect either T_i or T_j . However, the value obtained by the next $\text{read}(Q)$ instruction of S is affected, since the result of only the latter of the two write instructions is preserved in the database. If there is no other $\text{write}(Q)$ instruction after I and J in S , then the order of I and J directly affects the final value of Q in the database state that results from schedule S .
- We say that I and J **conflict** if they are operations by different transactions on the same data item, and at least one of these instructions is a write operation.


Conflicting Instructions

T_1	T_2
read (A) write (A)	read (A) write (A)
read (B) write (B)	read (B) write (B)

- ❑ write(A) instruction of T_1 conflicts with the read(A) instruction of T_2 .
- ❑ write(A) instruction of T_2 does not conflict with the read(B) instruction of T_1 because the two instructions access different data items.

Conflicting Instructions

- ❑ Let I and J be consecutive instructions of a schedule S .
- ❑ If I and J are instructions of different transactions and I and J do not conflict, then we can swap the order of I and J to produce a new schedule S' .
- ❑ S is equivalent to S' , since all instructions appear in the same order in both schedules except for I and J , whose order does not matter.



Conflicting Instructions

□ Continue to swap nonconflicting instructions:

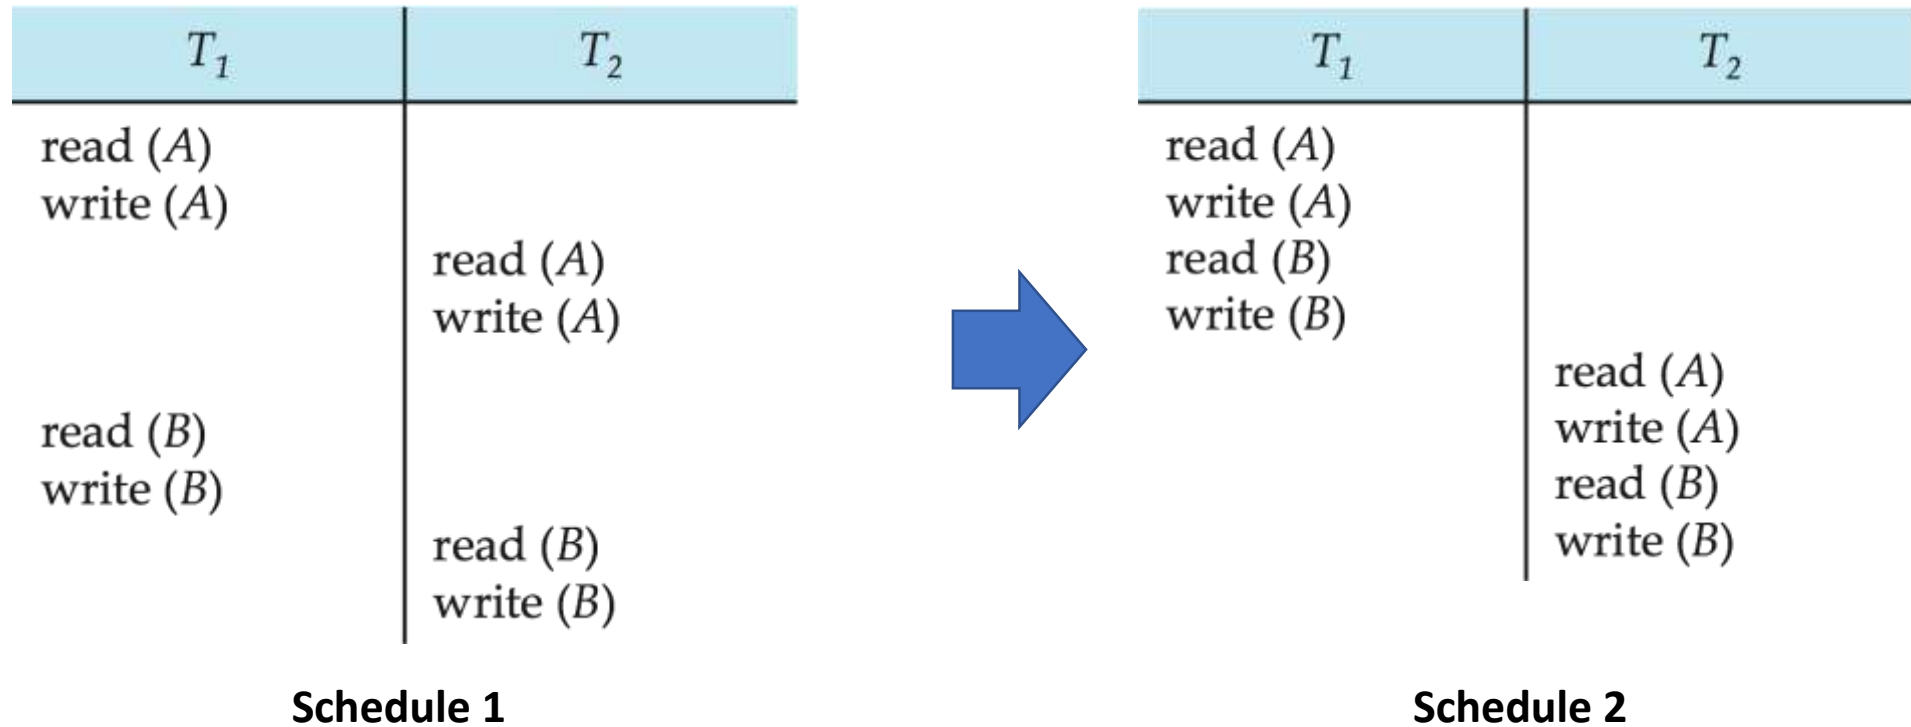
- Swap read(B) instruction of T_1 with read(A) instruction of T_2 .
- Swap write(B) instruction of T_1 with write(A) instruction of T_2 .
- Swap write(B) instruction of T_1 with read(A) instruction of T_2 .

T_1	T_2
read (A)	
write (A)	
read (B)	
write (B)	
	read (A)
	write (A)
	read (B)
	write (B)

Conflict Serializability

- ❑ If a schedule S can be transformed into a schedule S' by a series of swaps of nonconflicting instructions, we say that S and S' are **conflict equivalent**.
- ❑ A schedule S is **conflict serializable** if it is conflict equivalent to a serial schedule.

Conflict Serializability



- ❑ Schedule 1 can be transformed into Schedule 2 - a serial schedule where T_2 follows T_1 , by a series of swaps of non-conflicting instructions.
- ❑ Schedule 1 is conflict serializable.

Conflict Serializability

❑ Example of a schedule that is not conflict serializable:

T_3	T_4
read (Q)	
	write (Q)
write (Q)	


❑ We are unable to swap instructions in the above schedule to obtain either the serial schedule $\langle T_3, T_4 \rangle$, or the serial schedule $\langle T_4, T_3 \rangle$.

Precedence Graph


- ❑ Consider a schedule S .
- ❑ We construct a directed graph, called precedence graph, from S .
- ❑ This graph consists of a pair $G = (V, E)$, where V is a set of vertices and E is a set of edges.
- ❑ Set of vertices consists of all the transactions participating in the schedule.
- ❑ Set of edges consists of all edges $T_i \rightarrow T_j$ for which one of three conditions holds:
 - T_i executes $\text{write}(Q)$ before T_j executes $\text{read}(Q)$
 - T_i executes $\text{read}(Q)$ before T_j executes $\text{write}(Q)$
 - T_i executes $\text{write}(Q)$ before T_j executes $\text{write}(Q)$
- ❑ If an edge $T_i \rightarrow T_j$ exists in precedence graph, then, in any serial schedule S' equivalent to S , T_i must appear before T_j .

Precedence Graph

T_1	T_2
read (A) write (A) read (B) write (B)	read (A) write (A) read (B) write (B)




T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit



Precedence Graph



T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit



Testing for Conflict Serializability

- ❑ If the precedence graph for schedule S has a cycle, then schedule S is not conflict serializable.
- ❑ If the graph contains no cycles, then the schedule S is conflict serializable.
- ❑ A schedule is conflict serializable if and only if its precedence graph is acyclic.


Serializability Order



- ❑ To find an equivalent serial schedule, start with a node that does not have any incoming edges, and then make sure that the node order for every edge is not violated.

Testing for Conflict Serializability

	Transaction T_1	Transaction T_2	Transaction T_3
Time ↓	read_item(X); write_item(X);	read_item(Z); read_item(Y); write_item(Y);	read_item(Y); read_item(Z);
	read_item(Y); write_item(Y);	read_item(X); write_item(X);	write_item(Y); write_item(Z);



Equivalent serial schedules

None


Reason

Cycle $X(T_1 \rightarrow T_2), Y(T_2 \rightarrow T_1)$

Cycle $X(T_1 \rightarrow T_2), YZ(T_2 \rightarrow T_3), Y(T_3 \rightarrow T_1)$

Testing for Conflict Serializability

	Transaction T_1	Transaction T_2	Transaction T_3
Time ↓	<code>read_item(X);</code> <code>write_item(X);</code> <code>read_item(Y);</code> <code>write_item(Y);</code>	 <code>read_item(Z);</code> <code>read_item(Y);</code> <code>write_item(Y);</code> <code>read_item(X);</code> <code>write_item(X);</code>	<code>read_item(Y);</code> <code>read_item(Z);</code> <code>write_item(Y);</code> <code>write_item(Z);</code>



Equivalent serial schedules

$T_3 \rightarrow T_1 \rightarrow T_2$

Conflict Serializability

T_1	T_2
read(A) $A := A - 50$ write(A)	
	read(B) $B := B - 10$ write(B)
read(B) $B := B + 50$ write(B)	
	read(A) $A := A + 10$ write(A)

Conflict Serializable??

Conflict Serializable?

$r1(X); r3(X); w1(X); r2(X); w3(X);$

$S1: r1(X); r2(Z); r1(Z); r3(X); r3(Y); w1(X); w3(Y); r2(Y); w2(Z); w2(Y);$

Recoverable Schedules

- ❑ Recovery is possible.
- ❑ Nonrecoverable schedules should not be permitted by the DBMS.
- ❑ No committed transaction ever needs to be rolled back.
- ❑ **Recoverable schedule** — If a transaction T_j reads a data item previously written by a transaction T_i , then the commit operation of T_i must appear before the commit operation of T_j .

Recoverable Schedules

❑ Following schedule is not recoverable if T_9 commits immediately after the read(A) operation.

T_8	T_9
read (A) write (A)	
	read (A) commit
read (B)	

❑ If T_8 should abort, T_9 would have read (and possibly shown to the user) an inconsistent database state.

❑ Hence, database must ensure that schedules are recoverable.

Cascading Rollback

- ❑ Cascading rollback may occur in some recoverable schedules.
 - Uncommitted transaction may need to be rolled back.
- ❑ A single transaction failure leads to a series of transaction rollbacks.
- ❑ Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable):

T_{10}	T_{11}	T_{12}
read (A) read (B) write (A)	read (A) write (A)	
abort		read (A)

- ❑ If T_{10} fails, T_{11} and T_{12} must also be rolled back.
- ❑ Can lead to the undoing of a significant amount of work.

Cascadeless Schedule

- ❑ For each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the read operation of T_j .
- ❑ Avoids cascading rollback.
- ❑ Every cascadeless schedule is also recoverable.
- ❑ It is desirable to restrict the schedules to those that are cascadeless.
- ❑ Example of a schedule that is not cascadeless:

T_{10}	T_{11}	T_{12}
read (A) read (B) write (A)	read (A) write (A)	
abort		read (A)

Concurrency Control

- ❑ A database must provide a mechanism that will ensure that all possible schedules are both:
 - Conflict serializable
 - Recoverable and preferably cascadeless
- ❑ A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency.
- ❑ Concurrency-control schemes tradeoff between the amount of concurrency they allow and the amount of overhead that they incur.
- ❑ Testing a schedule for serializability *after* it has executed is a little too late!
 - Tests for serializability help us understand why a concurrency control protocol is correct.
- ❑ **Goal** – to develop concurrency control protocols that will assure serializability.

Concurrency Control

- ❑ Serializable schedule gives benefit of concurrent execution.
 - Without giving up any correctness
- ❑ DBMS enforces protocols
 - Set of rules to ensure serializability

View Equivalent

- ❑ Let S and S' be two schedules with the same set of transactions and include the same operations of those transactions.
- ❑ S and S' are **view equivalent** if the following three conditions are met for each data item Q :
 - If in schedule S , transaction T_i reads the initial value of Q , then in schedule S' also transaction T_i must read the initial value of Q .
 - If in schedule S transaction T_i executes $\text{read}(Q)$, and that value was produced by transaction T_j (if any), then in schedule S' also transaction T_i must read the value of Q that was produced by the same $\text{write}(Q)$ operation of transaction T_j .
 - The transaction (if any) that performs the final $\text{write}(Q)$ operation in schedule S must also perform the final $\text{write}(Q)$ operation in schedule S' .

View Serializability

- ❑ A schedule S is **view serializable** if it is view equivalent to a serial schedule.
- ❑ Any conflict serializable schedule is also view serializable but not vice versa.
- ❑ Below is a schedule which is view-serializable but not conflict serializable.

T_{27}	T_{28}	T_{29}
read (Q)	write (Q)	
write (Q)		write (Q)

- ❑ Every view serializable schedule that is not conflict serializable has **blind writes**.
- ❑ Problem of checking if a schedule is view serializable falls in the class of NP-hard.
 - Finding an efficient polynomial time algorithm for this problem is highly unlikely.

Neither Conflict nor View Serializable

- ❑ Schedule below produces the same outcome as the serial schedule $\langle T_1, T_5 \rangle$, yet is not conflict equivalent or view equivalent to a serial schedule.

T_1	T_5
read (A) $A := A - 50$ write (A)	
	read (B) $B := B - 10$ write (B)
read (B) $B := B + 50$ write (B)	
	read (A) $A := A + 10$ write (A)